

(by color or style of lines and markers). When asked to provide a brief explanation or description, try to make your answers concise, but do not omit anything you believe is important.

When submitting code, please make sure it's reasonably documented, and describe succinctly in the written component of the solution what is done in each py-file.

1 Regression

In this set of problems we will look at the regression problem and the maximum likelihood (ML) approach, with the goal to understand a bit better some of their properties.

We will start with simple linear regression, defined by

$$\hat{y} = f(\mathbf{x}; \mathbf{w}) = \mathbf{w} \cdot \mathbf{x} = \sum_{j=0}^d w_j x_j. \quad (1)$$

It is assumed in (1) that we have augmented the inputs $\mathbf{x} \in \mathbb{R}^d$ by adding 1 as the “zeroth” dimension, $x_0 \equiv 1$. This assumption is valid for the remainder of this problem set, unless stated otherwise.

Let $\hat{\mathbf{w}}$ be the linear regression parameters estimated using the least squares procedure from data $\mathbf{x}_1, \dots, \mathbf{x}_N$, and \mathbf{w}^* be the *best* possible linear regression parameters for the underlying $p(\mathbf{x}, y)$.

It was shown in the lectures that prediction errors made by $\hat{\mathbf{w}}$ are uncorrelated with the training $\{\mathbf{x}_i\}$, and a more general statement was made but not proven: that the prediction errors are in fact uncorrelated with values of any linear function of \mathbf{x} computed on the training inputs.

Later, we made use of a similar fact regarding the prediction errors made by \mathbf{w}^* . Here we will prove this more general fact. First, to make things precise, let us recall the definition of correlation between two (continuous) random variables U and V . Let $p_{u,v}$ be their joint probability density, and p_u and p_v the corresponding marginal densities. Then, the correlation (coefficient) between U and V is defined as

$$\rho(U, V) \triangleq \frac{E_{p_{u,v}} [(U - E_{p_u} [U]) (V - E_{p_v} [V])]}{\sqrt{\text{var}(U) \text{var}(V)}}, \quad (2)$$

where $\text{var}(U)$ and $\text{var}(V)$ are the variances of U and V under their respective marginal distributions. Intuitively, correlation measures how well one variable is linearly predictable from the other.

We will now prove a fact from which it follows that prediction errors of \mathbf{w}^* are uncorrelated with any linear function of \mathbf{x} :

Problem 1 [30 points]

Show that for any $\mathbf{a} \in \mathbb{R}^{d+1}$,

$$E_{p(\mathbf{x}, y)} [(y - \mathbf{w}^{*T} \mathbf{x}) \mathbf{a}^T \mathbf{x}] = 0. \quad (3)$$

End of problem 1

Advice: You may want to write explicitly what the definition of \mathbf{w}^ as the best linear regression under $p(\mathbf{x}, y)$ implies, in terms of minimizing the expected loss (which, to remind you, is an integral w.r.t. \mathbf{x} and y) similarly to how we treated the empirical loss in the case of $\hat{\mathbf{w}}$.*

Problem 2 [5 points]

Explain (succinctly but precisely) how the original statement, regarding zero correlation between the prediction errors made by $\hat{\mathbf{w}}$ estimated by least squares with any linear function of the training $\{\mathbf{x}_i\}$, follows from (3).

End of problem 2

Next, we consider the effect of *scaling* the input in the regression problem. This becomes relevant, for instance, in polynomial regression with high order models, to prevent very large or very small values and the resulting potential for numerical instability (think about the case of $x = 0.01$ or $x = 1000$ for 10-order model...) It seems that a good solution could be to multiply each column of the design matrix \mathbf{X} by a suitable number so that the range of that column (i.e. of the corresponding dimension in the regression input space) be fixed, say, to $[-1, 1]$.

Suppose that the data are scaled by multiplying the j -th dimension of the input by a non-zero number c_j . We will denote a single normalized data point by $\tilde{\mathbf{x}} \triangleq [1, c_1 x_1, \dots, c_d x_d]^T$ and the resulting design matrix by $\tilde{\mathbf{X}}$.

Problem 3 [25 points]

Let $\hat{\mathbf{w}}$ be the least squares estimate of the regression parameters from the unscaled \mathbf{X} , and let $\hat{\tilde{\mathbf{w}}}$ be the solution obtained from the scaled $\tilde{\mathbf{X}}$. Show that the scaling does not change optimality, in the sense that $\hat{\mathbf{w}} \cdot \mathbf{x} = \hat{\tilde{\mathbf{w}}} \cdot \tilde{\mathbf{x}}$.

End of problem 3

Advice: You may find it helpful to express the scaling as a linear operator, yielding a matrix-product expression, and to equip yourself with a matrix reference, such as the Matrix Cookbook (look it up if you are not already familiar with it!)

2 Asymmetric loss

In this section we will consider a regression model associated with a different loss function.

In class we have discussed the squared loss


$$\ell(\hat{y}, y) = (\hat{y} - y)^2, \quad (4)$$

which is *symmetric* – the sign of the prediction error in \hat{y} with respect to ground truth y doesn't matter. In some applications, this may not be a reasonable loss. For instance, consider the problem of estimating pressure necessary to put in the tire of a truck as a function of some parameters of the road, environment, vehicle etc. Putting pressure which is too low could reduce performance; if it's extremely low it can of course be dangerous. However, putting too much pressure may quickly result in catastrophic damage to the vehicle.

In this situation, we would like to penalize for positive errors differently (more severely) than for the negative errors of the magnitude. We can express this in the following, *asymmetric* squared loss:

$$\ell_{\alpha}(\hat{y}, y) = \begin{cases} (\hat{y} - y)^2 & \text{if } \hat{y} \leq y, \\ \alpha (\hat{y} - y)^2 & \text{if } \hat{y} \geq y. \end{cases} \quad (5)$$

The coefficient α determines how much more we worry about over-predicting y than about under-predicting. We will assume that $\alpha \geq 1$. See figure for an illustration.



Plot of asymmetric loss ℓ_{10} ; blue portion corresponds to negative errors (under-prediction), red to positive errors.

For the following problems, we have provided you with three data sets: training, validation and test, drawn from the same (unknown to you) $p(\mathbf{x}, y)$. The data are in a HDF5 file `data.h5`.

You will use training set to fit different models; validation to select between models, and test set to report the accuracy with the model that you end up choosing. That means that the test set will be only used once, to compute and report test loss with the model of choice.

For your convenience, we provide most of the code you will need in the experiments below, with a few missing pieces that you will need to fill in. If you are new to Python, don't worry too much about efficiency and elegance of your code, although those are of course good goals eventually; for now just make sure it works correctly and is readable.

The code “skeleton”, with quite a bit of flesh on it, is provided as a Jupyter notebook `PS1-asymloss-notebook.ipynb`, along with an auxiliary file `utils.py` containing a couple of functions you will need; take time to read and understand the provided code, especially if you are new to Python.

Please submit the notebook with all the code filled in and with the output of your experiments, after the name change as stated in the preamble of this problem set. Please write your comments, explanations and interpretation of the results as requested in the problems below, in the PDF solution file.

Problem 4 [15 points]

Fit linear, quadratic and cubic models to the training data under the standard (symmetric) squared loss (4). Visualize the fit under the three models. To make it easier to parse your figures, let's agree on the color scheme: the plots will be blue (linear), green (quadratic) and red (cubic) solid lines.

Evaluate, for each of the three models, the (symmetric) squared loss and the log-likelihood of the training data under the Gaussian noise assumption.

Note: you will need to derive ML estimate for the Gaussian noise variance σ^2 , in addition to the ML estimate for the model parameters \mathbf{w} ; write the missing code snippet for this, and explain in the solution file how you derived it.

Finally, compute and report (symmetric) squared loss for each of the three models on the validation set. Comment on the relative performance of the models, and on the gap, if any, between validation and training loss values. Select one of the three models based on this evaluation, and explain your selection. Let's call the model you end up choosing here model A.

End of problem 4

Problem 5 [25 points]

Now we want to repeat the experiment above but under the asymmetric loss function(5). Since there is no closed form solution, we will need to rely on gradient descent. Derive the expression for the gradient of the loss (5) with respect to the parameter vector \mathbf{w} , for a general form of linear regression (i.e., with inputs represented by a feature vector $\phi(\mathbf{x})$, which could capture linear, quadratic or cubic model).

Complete the provided code for gradient descent using the derivation for the gradient, and run it to fit linear, quadratic and cubic models under asymmetric squared loss. Visualize the resulting three models (use same color scheme as before, but dashed lines) and report the asymmetric squared loss values and the log-likelihood of the model *under the standard Gaussian noise model* on training data, and the asymmetric squared loss on validation data. Compare these numbers to those you obtained with symmetric loss, and comment on the relative performance of the three models, as well as the performance of the equal degree models across the two loss functions.

Next, using the outcome of your experiment, select a model among the three (linear/quadratic/cubic) and explain your selection. Let's call the resulting model B.

Evaluate model A (selected with symmetric loss) and model B on test set, and interpret your findings; which model do you conclude to be a better choice based on this entire experiment? Choice of evaluation metrics is up to you, and you should motivate it succinctly but clearly.

End of problem 5

Advice: You may want to start by reviewing the derivation of the gradient of squared loss in Lecture 2, and extend it to asymmetric loss.